

# Classification

(supervised learning)

JJ Valletta

April 7, 2015

[www.exeter.ac.uk/as/rdp/](http://www.exeter.ac.uk/as/rdp/)

# Overview

- The classification task
- Over/under-fitting a predictive model
- $k$ -nearest neighbour ( $k$ NN)
- Decision trees
- Random forests
- Support vector machines (SVM)

# Overview

- The classification task
  - Over/under-fitting a predictive model
  - $k$ -nearest neighbour ( $k$ NN)
  - Decision trees
  - Random forests
  - Support vector machines (SVM)

# Overview

- The classification task
- Over/under-fitting a predictive model
- $k$ -nearest neighbour ( $k$ NN)
- Decision trees
- Random forests
- Support vector machines (SVM)

# Overview

- The classification task
- Over/under-fitting a predictive model
- $k$ -nearest neighbour ( $k$ NN)
- Decision trees
- Random forests
- Support vector machines (SVM)

# Overview

- The classification task
- Over/under-fitting a predictive model
- $k$ -nearest neighbour ( $k$ NN)
- Decision trees
- Random forests
- Support vector machines (SVM)

# Overview

- The classification task
- Over/under-fitting a predictive model
- $k$ -nearest neighbour ( $k$ NN)
- Decision trees
- Random forests
- Support vector machines (SVM)

# Overview

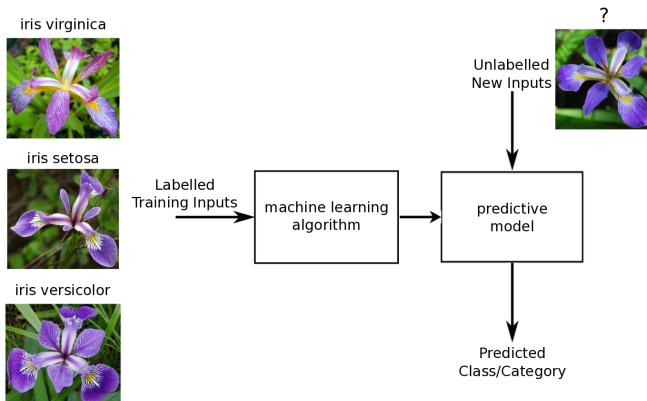
- The classification task
- Over/under-fitting a predictive model
- $k$ -nearest neighbour ( $k$ NN)
- Decision trees
- Random forests
- Support vector machines (SVM)



# Classification

## Definition

Assigning a set of new observations to a predefined category/class, using a predictive model trained on observations whose category/class is known



**Note:** Better data/features will almost always beat better algorithms

# Where are classifiers used?

**Medical imaging:** is tumour benign or cancerous

**Gene expression:** use “signature” to classify patient as having (or not) a condition

**Computer vision:** detect and track a moving object

**Biogeography:** classify land cover using remote sensing imagery

**Speech recognition:** translate audio signals into written text


**Biometric authentication:** identify a person using some personal characteristic e.g fingerprint or DNA

**Epidemiology:** given a set of risk factors what is the chance of patient suffering from a certain condition


# Over/under-fitting (bias-variance tradeoff)

- How well should we fit the training data to get good generalisation?
- Driving training error to zero is not a good idea
- **Bias** caused by a too rigid model leads to *underfitting*
- **Variance** caused by a too flexible model leads to *overfitting*
- **Occam's Razor**, pick the simplest model that explains your data, *parsimony*


1 degree polynomial (underfitting)




3 degree polynomial (parsimonious)



12 degree polynomial (overfitting)



# Iris dataset



## $k$ -nearest neighbour ( $k$ NN)

```
library(class)
fit <- knn(train, test, cl, k)
# train - training dataset (matrix or data frame)
# test - testing dataset (matrix or data frame)
# cl - corresponding label of the training dataset (factor)
# k - number of neighbours
```

- 1 Calculate distance between test point and every training data point
- 2 Find the  $k$  training points closest to test point
- 3 Assign test point the majority vote of their class label

## $k$ -nearest neighbour ( $k$ NN)

```
library(class)
fit <- knn(train, test, cl, k)
# train - training dataset (matrix or data frame)
# test - testing dataset (matrix or data frame)
# cl - corresponding label of the training dataset (factor)
# k - number of neighbours
```

- 1 Calculate distance between test point and every training data point
- 2 Find the  $k$  training points closest to test point
- 3 Assign test point the majority vote of their class label

## $k$ -nearest neighbour ( $k$ NN)

```
library(class)
fit <- knn(train, test, cl, k)
# train - training dataset (matrix or data frame)
# test - testing dataset (matrix or data frame)
# cl - corresponding label of the training dataset (factor)
# k - number of neighbours
```

- 1 Calculate distance between test point and every training data point
- 2 Find the  $k$  training points closest to test point
- 3 Assign test point the majority vote of their class label


## $k$ -nearest neighbour ( $k$ NN)

```
library(class)
fit <- knn(train, test, cl, k)
# train - training dataset (matrix or data frame)
# test - testing dataset (matrix or data frame)
# cl - corresponding label of the training dataset (factor)
# k - number of neighbours
```


- 1 Calculate distance between test point and every training data point
- 2 Find the  $k$  training points closest to test point
- 3 Assign test point the majority vote of their class label




# $k$ -nearest neighbour



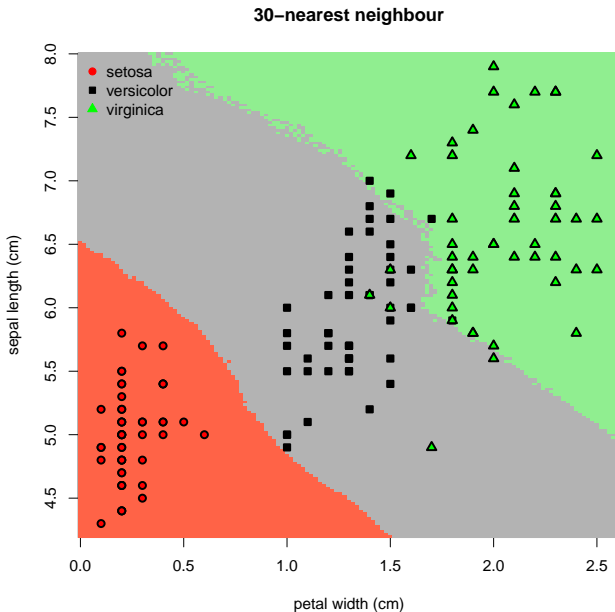
# $k$ -nearest neighbour



# $k$ -nearest neighbour



# $k$ -nearest neighbour



# $k$ -nearest neighbour

## Pros

- Simple and intuitive
- Works for multi-class problems
- Non-linear decision boundaries
- $k$  easily tuned by cross-validation

## Cons

- Can be computationally expensive as for every test point distance to every training data points needs to be computed (the model is actually the whole training dataset)
- Defining nearest by a distance metric can be ambiguous (for e.g when you have categorical predictors)

# Decision trees

```
library(tree)
fit <- tree(formula, data)
# OR
library(rpart)
fit <- rpart(formula, data)
# formula - an R formula expression e.g y ~ x1+x2
# data - data frame
```

- ① Divide data into left-right (yes-no) by an axis parallel split using *one* predictor
- ② The best split is found by maximising information gain/lowering entropy
- ③ Repeat 1 to 2 until all data is correctly classified or some stopping rule reached

# Decision trees

```
library(tree)
fit <- tree(formula, data)
# OR
library(rpart)
fit <- rpart(formula, data)
# formula - an R formula expression e.g y ~ x1+x2
# data - data frame
```

- 1 Divide data into left-right (yes-no) by an axis parallel split using *one* predictor
- 2 The best split is found by maximising information gain/lowering entropy
- 3 Repeat 1 to 2 until all data is correctly classified or some stopping rule reached

# Decision trees

```
library(tree)
fit <- tree(formula, data)
# OR
library(rpart)
fit <- rpart(formula, data)
# formula - an R formula expression e.g y ~ x1+x2
# data - data frame
```

- 1 Divide data into left-right (yes-no) by an axis parallel split using *one* predictor
- 2 The best split is found by maximising information gain/lowering entropy
- 3 Repeat 1 to 2 until all data is correctly classified or some stopping rule reached




# Decision trees

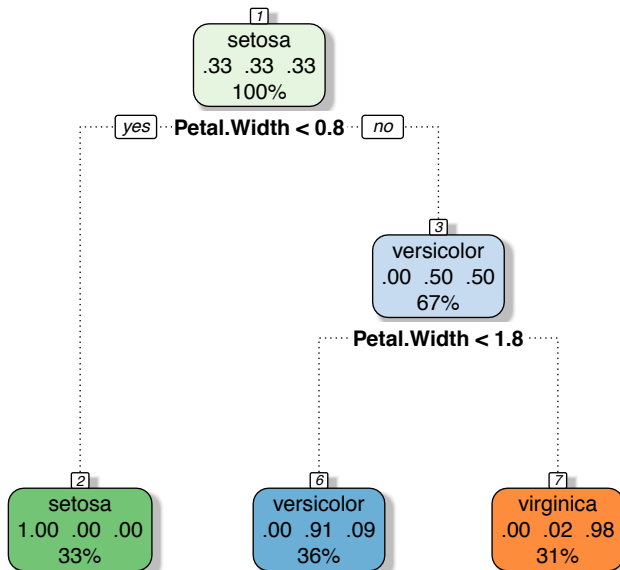
```
library(tree)
fit <- tree(formula, data)
# OR
library(rpart)
fit <- rpart(formula, data)
# formula - an R formula expression e.g y ~ x1+x2
# data - data frame
```

- 1 Divide data into left-right (yes-no) by an axis parallel split using *one* predictor
- 2 The best split is found by maximising information gain/lowering entropy
- 3 Repeat 1 to 2 until all data is correctly classified or some stopping rule reached


## Decision trees - The Model



# Decision trees - The Model



# Decision boundaries



# Decision trees

## Pros

- Model is very interpretable and can be directly used to generate rules
- Computationally inexpensive to train and evaluate
- Handle both categorical and continuous data
- Robust to outliers

## Cons

- Can easily overfit the data
- Predictive accuracy can be poor
- Small changes to training data may lead to a completely different tree

# Random forests

- Decision trees are intuitive but suffer from overfitting which significantly affect their predictive accuracy
- *Pruning*, to “trim” the tree back, help reduce this overfit
- *Ensemble* methods such as Random Forests are a better alternative
- **Rationale:** Instead of one tree, grow a *forest*, where every bushy tree (no pruning) is a bit different, then average predictions over all trees



# Random forests

```
library(randomForest)
fit <- randomForest(formula, data, ntree, mtry)
# formula - an R formula expression e.g y ~ x1+x2
# data - data frame
# ntree - number of trees in forest
# mtry - number of predictors randomly sampled as candidates at
        each split (default is sqrt(num of covariates))
```

- 1 Grow *ntree* bushy trees (no pruning) which are de-correlated from each other
- 2 Forest randomness is induced by:
  - *bootstrap* (Bernoulli resampling) - a subset of  $n$  data is randomly sampled with replacement
  - *mtry* - number of predictors randomly sampled at each split
- 3 Average predictions from all *ntree* trees

# Random forests

```
library(randomForest)
fit <- randomForest(formula, data, ntree, mtry)
# formula - an R formula expression e.g y ~ x1+x2
# data - data frame
# ntree - number of trees in forest
# mtry - number of predictors randomly sampled as candidates at
        each split (default is sqrt(num of covariates))
```

- 1 Grow *ntree* bushy trees (no pruning) which are de-correlated from each other
- 2 Forest randomness is induced by:
  - Randomly sampling *mtry* predictors from the full set of predictors
  - Randomly sampling *ntree* trees from the data
- 3 Average predictions from all *ntree* trees



# Random forests

```
library(randomForest)
fit <- randomForest(formula, data, ntree, mtry)
# formula - an R formula expression e.g y ~ x1+x2
# data - data frame
# ntree - number of trees in forest
# mtry - number of predictors randomly sampled as candidates at
#       each split (default is sqrt(num of covariates))
```


- 1 Grow *ntree* bushy trees (no pruning) which are de-correlated from each other
- 2 Forest randomness is induced by:
  - Bagging (**B**ootstrap **AGG**regat**ING**), each tree is trained on a subset of the data randomly sampled with replacement
  - For every tree split consider only *mtry* predictors as candidates for that split
- 3 Average predictions from all *ntree* trees

# Random forests


```
library(randomForest)
fit <- randomForest(formula, data, ntree, mtry)
# formula - an R formula expression e.g y ~ x1+x2
# data - data frame
# ntree - number of trees in forest
# mtry - number of predictors randomly sampled as candidates at
        each split (default is sqrt(num of covariates))
```

- 1 Grow *ntree* bushy trees (no pruning) which are de-correlated from each other
- 2 Forest randomness is induced by:
  - Bagging (**B**ootstrap **AGG**regat**ING**), each tree is trained on a subset of the data randomly sampled with replacement
  - For every tree split consider only *mtry* predictors as candidates for that split
- 3 Average predictions from all *ntree* trees



# De-correlated bushy trees in the forest



# De-correlated bushy trees in the forest



# Variable importance



# Random forests

## Pros


- State-of-the-art predictive accuracy
- Can handle thousands of both categorical and continuous predictors without variable deletion
- Robust to outliers
- Estimates the importance of every predictor
- Out-of-bag error (unbiased estimate of test error for every tree built)
- Can cope with unbalanced datasets by setting class weights

## Cons

- Harder to interpret than plain decision trees


# Support vector machines (SVM)

Which is the best separating line?



# Support vector machines (SVM)


Which is the best separating line?






# Support vector machines (SVM)

Which is the best separating line?




# Support vector machines (SVM)

Which is the best separating line?



# Support vector machines


**Rationale:** Maximise the *margin*, the distance to separating hyperplane



# Support vector machines


If it's something weird and it don't look good. Who ya gonna call?  
Ghostbusters?

*Close*, we need alternate dimensions to make them linearly separable




# Support vector machines

- Map data to a higher dimensional space where classes are linearly separable (artificially increasing number of predictors)
- $(x_1, x_2) \rightarrow (1, x_1, x_2, x_1x_2, x_1^2, x_2^2)$
- Hyperplane in *new* space is a conic section in *original* space




# Support vector machines: Simple example from 1D to 2D

1D (original) data is not linearly separable




# Support vector machines: Simple example from 1D to 2D

1D (original) data is not linearly separable



2D (transformed) data is now linearly separable



# Support vector machines - The kernel trick

- So our solution is to blow up the dimensions?
- But what about the “curse of dimensionality”?
- Very computationally expensive to work in high dimensions



# Support vector machines - The kernel trick

- So our solution is to blow up the dimensions?
- But what about the “curse of dimensionality”?
- Very computationally expensive to work in high dimensions

**Kernel trick** to the rescue! We work in an *implicit* feature space, that is, data is never explicitly computed in higher dimensions. Instead we only need to compute the pair-wise inner product.


**p.s** Kernel methods are mathematically intricate and beyond the scope of this workshop so I will stop here

# Support vector machines


```
library(e1071)
fit <- svm(formula, data, type, kernel)
# formula - an R formula expression e.g y ~ x1+x2
# data - data frame
# type - whether the problem is classification or regression
# kernel - the kernel function e.g "linear" or "radial basis"
```

- 1 Choose carefully a kernel function
- 2 Run optimiser to find maximum margin

# Support vector machines



# Support vector machines



# Support vector machines

## Pros

- State-of-the-art predictive accuracy
- Less prone to overfitting
- Only need to store the support vectors for the predictive model
- Picking the right kernel gives you flexibility and predictive power
- Global optimum guaranteed

## Cons

- Unintuitive/a black box
- Cannot visualise the feature space