

# Clustering

(unsupervised learning)

John Joseph Valletta

University of Exeter, Penryn Campus, UK

1st-2nd June 2015

[www.exeter.ac.uk/as/rdp/](http://www.exeter.ac.uk/as/rdp/)

# Overview

- What is clustering?
- Major types of clustering methods
- $k$ -means clustering
- Agglomerative hierarchical clustering
- Gaussian mixture models
- How do we determine the correct number of clusters?

# Overview

- What is clustering?
- Major types of clustering methods
- $k$ -means clustering
- Agglomerative hierarchical clustering
- Gaussian mixture models
- How do we determine the correct number of clusters?

# Overview

- What is clustering?
- Major types of clustering methods
  - $k$ -means clustering
  - Agglomerative hierarchical clustering
  - Gaussian mixture models
  - How do we determine the correct number of clusters?

# Overview

- What is clustering?
- Major types of clustering methods
- $k$ -means clustering
- Agglomerative hierarchical clustering
- Gaussian mixture models
- How do we determine the correct number of clusters?

# Overview

- What is clustering?
- Major types of clustering methods
- $k$ -means clustering
- Agglomerative hierarchical clustering
- Gaussian mixture models
- How do we determine the correct number of clusters?

# Overview

- What is clustering?
- Major types of clustering methods
- $k$ -means clustering
- Agglomerative hierarchical clustering
- Gaussian mixture models
- How do we determine the correct number of clusters?

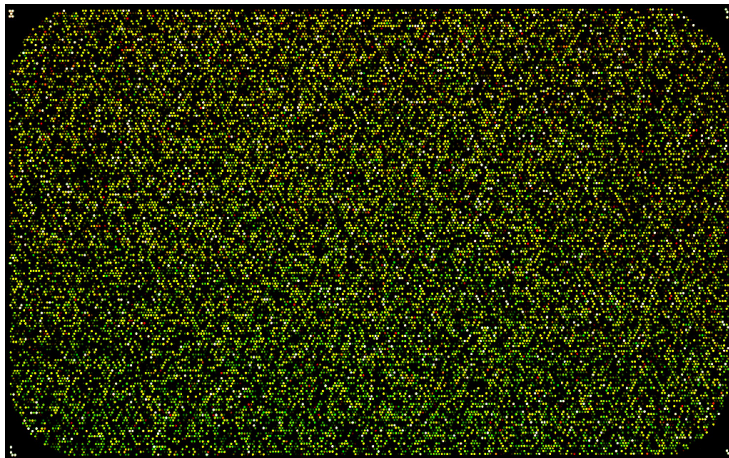
# Overview

- What is clustering?
- Major types of clustering methods
- $k$ -means clustering
- Agglomerative hierarchical clustering
- Gaussian mixture models
- How do we determine the correct number of clusters?



# What's the problem?

Well, it's driven by your data, e.g which of these genes are co-regulated?



We want to find some underlying structure in the data → **Clustering**

# What's the problem?

Gene expression: discovering co-regulated genes

Biological systematics: finding organisms sharing similar attributes

Computer vision: segmenting a digital image for object recognition

Epidemiology: identifying geographical clusters of diseases

Medical imaging: differentiating between tissues

Mathematical chemistry: grouping compounds by topological indices

**Clustering is particularly useful in applications where labelling the data is very time consuming/expensive**

# What is clustering?

## Formal definition

Identifying homogeneous and well separated groups of data points (features) by some similarity measure

# What is clustering?

## Formal definition

Identifying homogeneous and well separated groups of data points (features) by some similarity measure

## Informal definition

The process of stereotyping your data

e.g. *these* are round(ish) faces, *these* are short(ish) people

# What is clustering?

## Formal definition

Identifying homogeneous and well separated groups of data points (features) by some similarity measure

## Informal definition

The process of stereotyping your data  
e.g. *these* are round(ish) faces, *these* are short(ish) people

## But how many groups?

An unsolved problem. Issue lies in the subjectivity of the word **similar** and its mathematical definition

Are they similar?



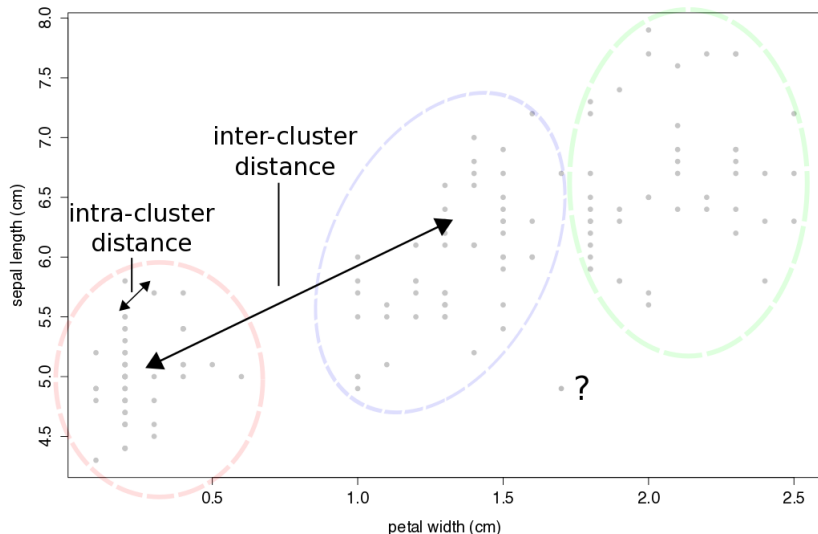
Are they similar?



# What are we after?

**Motivation:** How is the data structured? Any outliers?

**Goals:** *High* intra-cluster similarity and *low* inter-cluster similarity





# Major types of clustering methods


Partitional: The data (feature) space is partitioned into  $k$  regions

Hierarchical: Iteratively merging small clusters into larger ones (*agglomerative*) or breaking large clusters into smaller ones (*divisive*)

Distribution-based: Fit  $k$  multivariate statistical distributions

# Major types of clustering methods

**Partitional:** The data (feature) space is partitioned into  $k$  regions




**Hierarchical:** Iteratively merging small clusters into larger ones (*agglomerative*) or breaking large clusters into smaller ones (*divisive*)

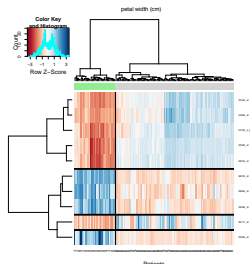
**Distribution-based:** Fit  $k$  multivariate statistical distributions

# Major types of clustering methods

**Partitional:** The data (feature) space is partitioned into  $k$  regions




**Hierarchical:** Iteratively merging small clusters into larger ones (*agglomerative*) or breaking large clusters into smaller ones (*divisive*)




**Distribution-based:** Fit  $k$  multivariate statistical distributions

# Major types of clustering methods


**Partitional:** The data (feature) space is partitioned into  $k$  regions



**Hierarchical:** Iteratively merging small clusters into larger ones (*agglomerative*) or breaking large clusters into smaller ones (*divisive*)

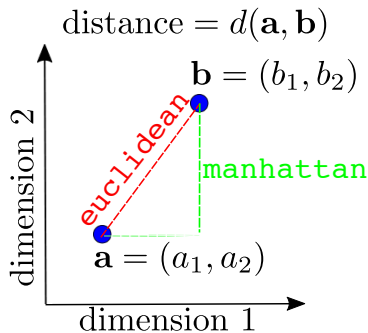


**Distribution-based:** Fit  $k$  multivariate statistical distributions



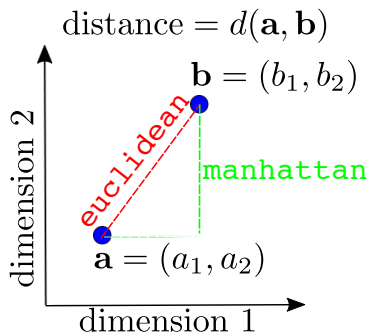
# Similarity measures

- A *distance* metric quantifies how *close* two data points are
- Several ways to define this distance which has a direct impact on the clustering result

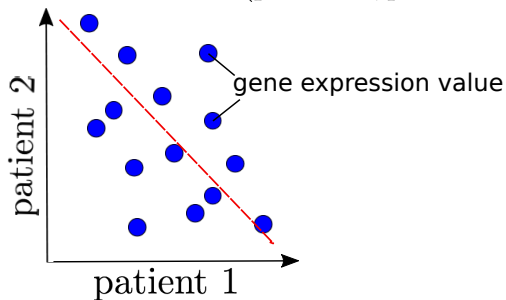


# Similarity measures

- A *distance* metric quantifies how *close* two data points are
- Several ways to define this distance which has a direct impact on the clustering result



distance =  $1 - \text{corr}(\text{patient 1}, \text{patient 2})$



# $k$ -means clustering

```
fit <- kmeans(x, centers)
# x - numeric matrix of data
# centers - no. of clusters k
```

- 1 Select  $k$  centroids at random
- 2 Calculate distance between centroids and each data point
- 3 Assign each data point to the closest centroid
- 4 Compute new centroids; the average of all data points in that cluster
- 5 Repeat steps 2 to 4 until data points remain in the same cluster or some maximum number of iterations reached

# $k$ -means clustering

```
fit <- kmeans(x, centers)
# x - numeric matrix of data
# centers - no. of clusters k
```

- 1 Select  $k$  centroids at random
- 2 Calculate distance between centroids and each data point
- 3 Assign each data point to the closest centroid
- 4 Compute new centroids; the average of all data points in that cluster
- 5 Repeat steps 2 to 4 until data points remain in the same cluster or some maximum number of iterations reached



# $k$ -means clustering

```
fit <- kmeans(x, centers)
# x - numeric matrix of data
# centers - no. of clusters k
```

- 1 Select  $k$  centroids at random
- 2 Calculate distance between centroids and each data point
- 3 Assign each data point to the closest centroid
- 4 Compute new centroids; the average of all data points in that cluster
- 5 Repeat steps 2 to 4 until data points remain in the same cluster or some maximum number of iterations reached

# $k$ -means clustering

```
fit <- kmeans(x, centers)
# x - numeric matrix of data
# centers - no. of clusters k
```

- 1 Select  $k$  centroids at random
- 2 Calculate distance between centroids and each data point
- 3 Assign each data point to the closest centroid
- 4 Compute new centroids; the average of all data points in that cluster
- 5 Repeat steps 2 to 4 until data points remain in the same cluster or some maximum number of iterations reached

# $k$ -means clustering

```
fit <- kmeans(x, centers)
# x - numeric matrix of data
# centers - no. of clusters k
```

- 1 Select  $k$  centroids at random
- 2 Calculate distance between centroids and each data point
- 3 Assign each data point to the closest centroid
- 4 Compute new centroids; the average of all data points in that cluster
- 5 Repeat steps 2 to 4 until data points remain in the same cluster or some maximum number of iterations reached

# $k$ -means clustering

```
fit <- kmeans(x, centers)
# x - numeric matrix of data
# centers - no. of clusters k
```

- 1 Select  $k$  centroids at random
- 2 Calculate distance between centroids and each data point
- 3 Assign each data point to the closest centroid
- 4 Compute new centroids; the average of all data points in that cluster
- 5 Repeat steps 2 to 4 until data points remain in the same cluster or some maximum number of iterations reached


# $k$ -means clustering

```
fit <- kmeans(x, centers)
# x - numeric matrix of data
# centers - no. of clusters k
```


- 1 Select  $k$  centroids at random
- 2 Calculate distance between centroids and each data point
- 3 Assign each data point to the closest centroid
- 4 Compute new centroids; the average of all data points in that cluster
- 5 Repeat steps 2 to 4 until data points remain in the same cluster or some maximum number of iterations reached

**Note:**  $k$ -means clustering should *only* be used with continuous data


# $k$ -means clustering



# $k$ -means clustering




# $k$ -means clustering

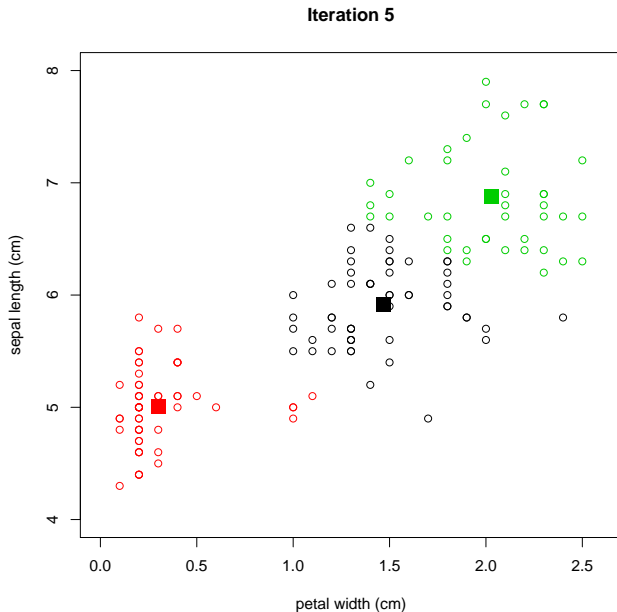




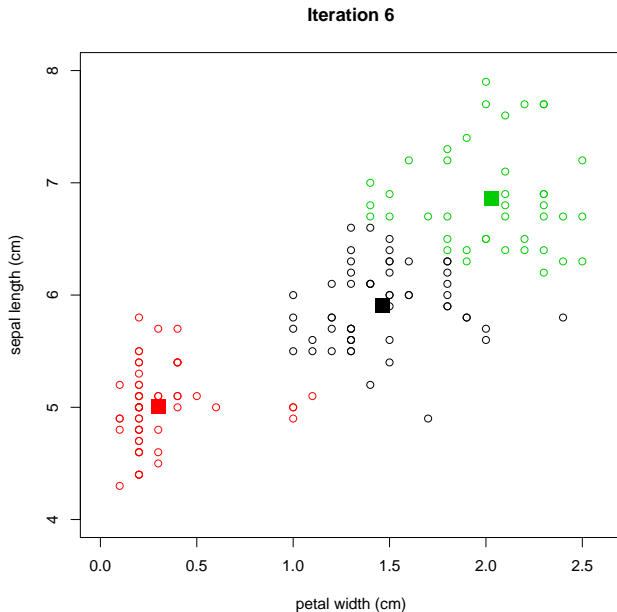
# $k$ -means clustering




# $k$ -means clustering



# $k$ -means clustering



# $k$ -means clustering



# $k$ -means clustering

## Pros

- Simple and intuitive
- Computationally inexpensive/fast

## Cons

- What is  $k$ ?
- Only applicable to continuous data where a mean is defined
- No guarantee of a global optimum solution

# Agglomerative hierarchical clustering

```
d <- dist(as.matrix(data), method)
# data - data frame
# method - distance method e.g "euclidean" or "manhattan"
fit <- hclust(d, method)
# method - linkage function e.g "complete" or "single"
```

- ① Assign each data point as its own cluster
- ② Compute distance between each cluster
- ③ Merge the closest pair into a single cluster
- ④ Repeat 2 to 3 until you're left with one cluster

# Agglomerative hierarchical clustering

```
d <- dist(as.matrix(data), method)
# data - data frame
# method - distance method e.g "euclidean" or "manhattan"
fit <- hclust(d, method)
# method - linkage function e.g "complete" or "single"
```

- 1 Assign each data point as its own cluster
- 2 Compute distance between each cluster
- 3 Merge the closest pair into a single cluster
- 4 Repeat 2 to 3 until you're left with one cluster

# Agglomerative hierarchical clustering

```
d <- dist(as.matrix(data), method)
# data - data frame
# method - distance method e.g "euclidean" or "manhattan"
fit <- hclust(d, method)
# method - linkage function e.g "complete" or "single"
```

- 1 Assign each data point as its own cluster
- 2 Compute distance between each cluster
- 3 Merge the closest pair into a single cluster
- 4 Repeat 2 to 3 until you're left with one cluster



# Agglomerative hierarchical clustering

```
d <- dist(as.matrix(data), method)
# data - data frame
# method - distance method e.g "euclidean" or "manhattan"
fit <- hclust(d, method)
# method - linkage function e.g "complete" or "single"
```

- 1 Assign each data point as its own cluster
- 2 Compute distance between each cluster
- 3 Merge the closest pair into a single cluster
- 4 Repeat 2 to 3 until you're left with one cluster

# Agglomerative hierarchical clustering

```
d <- dist(as.matrix(data), method)
# data - data frame
# method - distance method e.g "euclidean" or "manhattan"
fit <- hclust(d, method)
# method - linkage function e.g "complete" or "single"
```

- 1 Assign each data point as its own cluster
- 2 Compute distance between each cluster
- 3 Merge the closest pair into a single cluster
- 4 Repeat 2 to 3 until you're left with one cluster

# Agglomerative hierarchical clustering

```
d <- dist(as.matrix(data), method)
# data - data frame
# method - distance method e.g "euclidean" or "manhattan"
fit <- hclust(d, method)
# method - linkage function e.g "complete" or "single"
```

- 1 Assign each data point as its own cluster
- 2 Compute distance between each cluster
- 3 Merge the closest pair into a single cluster
- 4 Repeat 2 to 3 until you're left with one cluster

**Note:** Step 3 is *key*, the distance method and linkage function dictate the final result

# Hierarchical clustering: Link method

How do we compute the inter-cluster distance? The *linkage function*

Centroid: mean of data points (same as in  $k$ -means)

Single: distance between closest pair of points


Complete: distance between furthest pair of points

Average: mean pairwise distance between all points

# Hierarchical clustering: Link method

How do we compute the inter-cluster distance? The *linkage function*

**Centroid:** mean of data points (same as in  $k$ -means)



Single: distance between closest pair of points

Complete: distance between furthest pair of points

Average: mean pairwise distance between all points

# Hierarchical clustering: Link method


How do we compute the inter-cluster distance? The *linkage function*

**Centroid:** mean of data points (same as in  $k$ -means)

**Single:** distance between closest pair of points

**Complete:** distance between furthest pair of points

**Average:** mean pairwise distance between all points



# Hierarchical clustering: Link method


How do we compute the inter-cluster distance? The *linkage function*

**Centroid:** mean of data points (same as in  $k$ -means)

**Single:** distance between closest pair of points

**Complete:** distance between furthest pair of points

**Average:** mean pairwise distance between all points



# Hierarchical clustering: Link method


How do we compute the inter-cluster distance? The *linkage function*

**Centroid:** mean of data points (same as in  $k$ -means)

**Single:** distance between closest pair of points

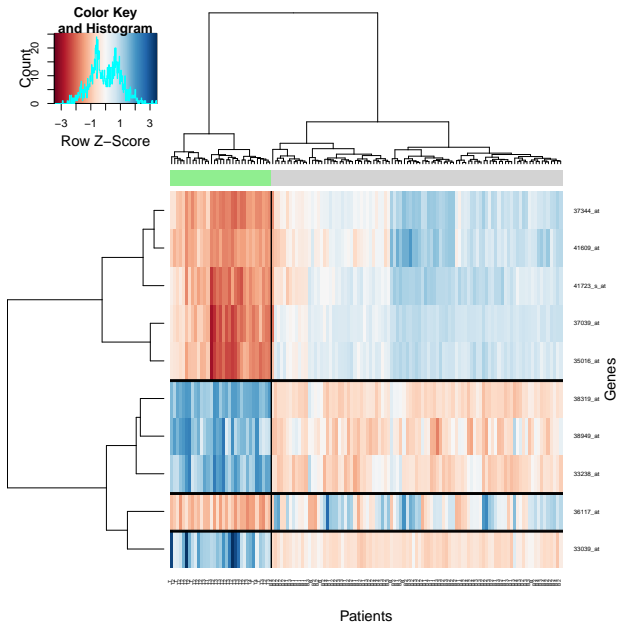
**Complete:** distance between furthest pair of points

**Average:** mean pairwise distance between all points





# Hierarchical clustering in gene expression studies



# Hierarchical clustering

## Pros

- No need to specify  $k$
- Results can be visualised nicely irrespective of number of dimensions

## Cons

- Can be computationally expensive
- Interpretation is subjective. Where should we draw the line (to separate clusters)?
- Choice of distance method and linkage function can significantly change the result

# Gaussian mixture models

```
library(mclust)
fit <- Mclust(data, G)
# data - data frame
# G - no. of Gaussians
```

## ① Fit $k$ multivariate Gaussian distributions

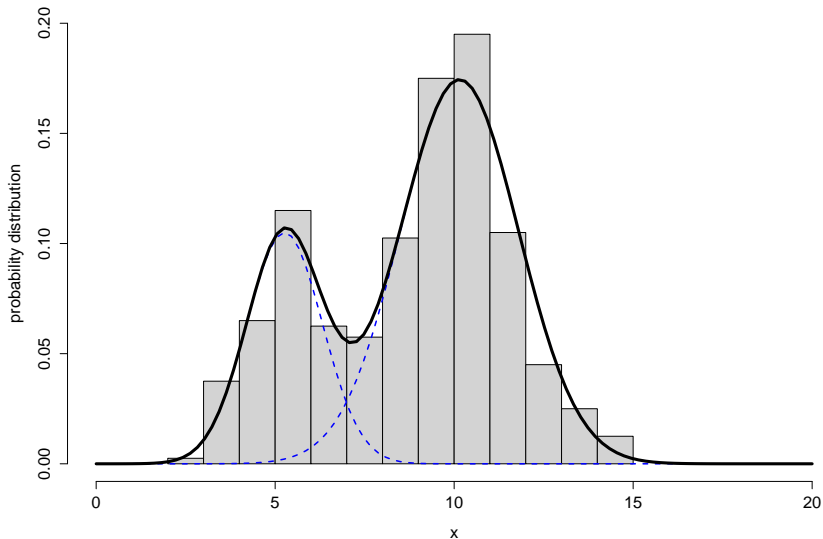
The Expectation-Maximisation (EM) algorithm is used to estimate the parameters  $\pi_i$  (mixing coefficients),  $\mu_i$  and  $\sigma_i$

$$p(x) = \sum_{i=1}^k \pi_i \mathcal{N}(x|\mu_i, \Sigma_i) \text{ and } \sum_{i=1}^k \pi_i = 1$$

Can be seen as a “soft” version of  $k$ -means because *every* point is part of *every* cluster but with varying levels of membership

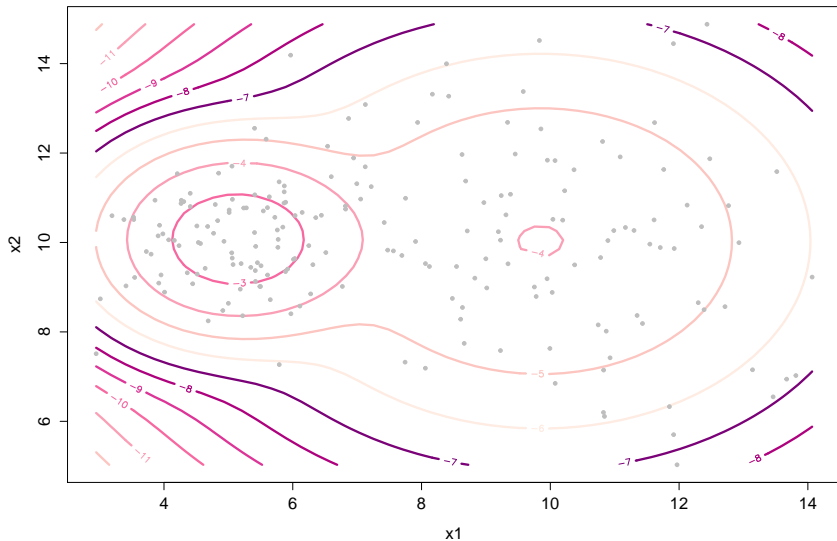
# Gaussian mixture models

Density Plot



# Gaussian mixture models

log Density Contour Plot



# Gaussian mixture models

## Pros

- Intuitive interpretation
- Computationally inexpensive

## Cons

- What is  $k$ ?
- Strong assumption on the data (normality)
- No guarantee of a global optimum solution

# How do we determine the correct number of clusters?



**Short answer:** you can't

Because data is unlabelled the correct number of  $k$  is ambiguous

However we can plot some indices as a function of  $k$  to help us evaluate cluster validity:

- Within and between clusters sum-of-squares distances
- Akaike Information Criterion (AIC) or Bayesian Information Criterion (BIC) when using distribution-based methods
- Silhouette plot

# How do we determine the correct number of clusters?





# How do we determine the correct number of clusters?

The NbClust package provides 30 different cluster validity metrics. A majority vote can be taken to deduce the appropriate number of clusters

```
library(NbClust)
NbClust(data, distance, method, min.nc, max.nc, index)
# data - data frame
# distance - similarity measure e.g "euclidean"
# method - clustering algorithm e.g "kmeans"
# min.nc - min number of clusters to consider
# max.nc - max number of clusters to consider
# index - which indices to compute, "all" computes all of them
```

## Note:

- Indices *only* give us a ballpark range for “correct” number of clusters
- Are they biologically relevant? Prior knowledge to the rescue
- E.g how many different phenotypes are you expecting?