

# Introduction to Machine Learning for the Life Sciences

## Lab Sessions

by

John Joseph Valletta  
University of Exeter, Penryn Campus, UK

1st-2nd June 2015

### Introduction

The aim of these lab sessions is to get you started with applying machine learning algorithms to various types of datasets. The focus is to showcase the versatility of these methods and give you a flavour of the sort of problems that can be tackled. We will not be spending time on the topic of *feature extraction*, that is, how to extract features which have the best predictive capabilities. This is usually the crux of any predictive modelling problem but which is domain-specific.

We will be using the programming language **R**, mainly because most life scientists are already familiar with it, but also because it is an open-source platform that gives you access to state-of-the-art algorithms. **Python** (free and extremely popular with computer scientists) and **MATLAB**® are other popular choices. You can also find faster implementations in **C/C++**.

These lab sessions are not intended to be an **R** tutorial. If you want to learn more about **R**, please refer to the myriad of free resources available on the web.

Installing **R**: <http://www.r-project.org/>

Installing RStudio: <http://www.rstudio.com/>

<p><b>It is imperative that you read and familiarise yourself with the documentation of any machine learning algorithm before using it!</b></p>
---

Everyone might have different packages available in their **R** installation, so before calling `library(Package Name)`, please run `install.packages("Package Name")`, where Package Name is the name of the package you need e.g `mclust`

# 1 Getting started using an artificial dataset

**Data:** artificial data from a 3D multivariate Gaussian distribution  
**Task:** cluster data points (unsupervised)  
**Method:**  $k$ -means and Gaussian mixture models

Before starting using real datasets, let us apply some of the algorithms we discussed to an artificial dataset that we can visualise and control. We will generate 5 sets of data from a 3-dimensional Gaussian distribution and visualise them.

```
1 library(MASS) # mvrnorm (multivariate normal)
2 library(rgl) # plot3d
3 N <- 100 # number of data points
4 sigma2 <- 1 # variance of data points
5 covMatrix <- sigma2*diag(3) # assume the same variance in all three directions
6 groupA <- mvrnorm(n=N, mu=c(3, 3, 9), Sigma=covMatrix)
7 groupB <- mvrnorm(n=N, mu=c(3, 9, 3), Sigma=covMatrix)
8 groupC <- mvrnorm(n=N, mu=c(9, 9, 9), Sigma=covMatrix)
9 groupD <- mvrnorm(n=N, mu=c(9, 3, 3), Sigma=covMatrix)
10 groupE <- mvrnorm(n=N, mu=c(6, 6, 6), Sigma=covMatrix)
11 xTrain <- rbind(groupA, groupB, groupC, groupD, groupE) # training dataset is 5*N x 3
12 plot3d(xTrain, col="grey", xlab="x", ylab="y", zlab="z")
```

Visually you can probably deduce that there are 5 different clusters. In practice however, it is hard to visualise high-dimensional data and we need to resort to other methods to deduce the number of clusters. The main objective of clustering is:

- High intra-cluster similarity (points in the same cluster should be similar)
- Low inter-cluster similarity (points in distinct clusters should be sufficiently different)

To guess the right number of clusters we can run  $k$ -means several times and plot the resultant intra- and inter-cluster sum-of-squares as a function of  $k$

```
1 # k-means, let k vary from 2 to 10
2 kRange <- seq(from=2, to=10, by=1)
3 intraClustSS <- rep(NA, length(kRange))
4 interClustSS <- rep(NA, length(kRange))
5 for (k in kRange){
6   fit <- kmeans(x=xTrain, centers=k)
7   intraClustSS[k-1] <- fit$tot.withinss # it's (k-1) because k starts from 2
8   interClustSS[k-1] <- fit$betweenss
9 }
10 minY <- min(c(intraClustSS, interClustSS))
11 maxY <- max(c(intraClustSS, interClustSS))
12 plot(kRange, interClustSS, type="o", pch=1, col="blue", lwd=3, lty=1, ylim=c(minY, maxY),
13      , xlab="k")
13 points(kRange, intraClustSS, type="o", pch=1, col="red", lwd=3)
14 legend("topleft", c("inter-cluster sum-of-squares", "intra-cluster sum-of-squares"), bty
15      ="n",
16      col=c("blue", "red"), pch=1, lwd=3, lty=1)
```

After  $k = 5$ , the sum-of-squares does not change that much, so it can be deduced that there are 5 distinct clusters in this dataset. Let us visualise the result:

```
1 fit <- kmeans(x=xTrain, centers=5)
2 plot3d(xTrain, col=fit$cluster, xlab="x", ylab="y", zlab="z")
```

**Question(s):**

1. What happens to the inter/intra-cluster sum-of-squares plot as we increase `sigma2`?

Let us perform a similar analysis but using a mixture of Gaussian models, where we will fit  $k$  Gaussians. This time, instead of the inter/intra-cluster sum-of-squares we will plot the BIC (Bayesian Information Criterion), a goodness-of-fit measure that penalises model complexity ( $k$ ). The `Mclust` function uses BIC to decide which model to retain. To speed up the process we will also tell `Mclust` that we want a Gaussian with equal variances in all directions (`modelName="EII"`).

```
1 library(mclust) # Gaussian mixture models
2 fit <- Mclust(data=xTrain, G=seq(10), modelName="EII")
3 summary(fit)
4 plot(fit$BIC, type="o", pch=1, col="blue", lwd=3, lty=1, xlab="k")
5 legend("topleft", "BIC (Bayesian Information Criterion)", bty="n",
6        col=c("blue", "red"), pch=1, lwd=3, lty=1)
7 plot(fit, what="density", lwd=2)
```

**Question(s):**

1. Repeat the above but without specifying `modelName`. Which model does it retain?
2. How does the BIC plot and the type of model retained changes when increasing `sigma2`?

## 2 Clustering gene expression data

<b>Data:</b>	gene expression ( <a href="#">data of T- and B-cell Acute Lymphocytic Leukemia</a> )
<b>Task:</b>	deduce the number of distinct phenotypes (unsupervised)
<b>Method:</b>	agglomerative hierarchical clustering

Agglomerative hierarchical clustering is a popular choice in gene expression studies. The data comes as a  $N \times D$  matrix, where  $N$  is the number of genes and  $D$  is the number of biological samples (e.g patients with different medical conditions or different tissue samples). The numerical value is a measure of how much the  $n$ th gene is expressed for the  $d$ th biological sample. There are usually two key questions that can be addressed with clustering:

1. Which biological samples exhibit a similar gene expression *signature*?
2. Which genes vary in a similar fashion?

Here we will use the popular acute lymphoblastic leukaemia (ALL) dataset. It is a 12625 (genes)  $\times$  128 (patients) microarray dataset. The patients were diagnosed with either a B- or T-cell acute lymphocytic leukaemia. Because we have access to the labels (type and stage of the disease e.g B2) we can easily assess how well the clustering algorithm is doing as we expect the B's and T's to cluster together. Note however that class labels are not always available and that the clustering algorithm does *not* know what the labels are. It is an unsupervised problem (we will turn it into a supervised problem later on).

The ALL dataset comes in the form of an `exprSet` object, a data structure used in [Bioconductor](#). Bioconductor is an ensemble of **R** packages specifically written to analyse high-throughput genomic data. Let us install Bioconductor (for those who have not got it installed already) and extract the  $N \times D$  gene expression matrix of interest <sup>1</sup>.

```
1 source("http://bioconductor.org/biocLite.R") # Install Bioconductor
2 biocLite() # Install core packages
3 biocLite("ALL") # Install the ALL package (takes a few secs)
4 library(ALL) # Load the ALL package
5 data(ALL) # Loads ALL dataset to workspace
6 xTrain <- exprs(ALL) # Extract the 12625 x 128 dataset
7 colnames(xTrain) <- ALL$BT # Replace patient ID by disease stage to assess clustering
```

For the purpose of this exercise, let us just focus on clustering the 128 patients:

```
1 distance <- dist(as.matrix(t(xTrain)), method="euclidean") # Distance between patients
2 # Perform agglomerative hierarchical clustering using the "complete" linkage function
3 fit <- hclust(distance, method="complete")
4 plot(fit, cex=0.5) # Visualise the resultant dendrogram
```

---

<sup>1</sup>Do not worry too much about the details, the objective is just to extract the data, but if your work involves this type of data it is recommended that you familiarise yourself with Bioconductor

**Question(s):**

1. Keep the distance method fixed and change the linkage method (e.g single, average). What happens to the dendrogram?
2. Keep the linkage method fixed and change the distance method (e.g manhattan, minkowski). What happens to the dendrogram?
3. Try using a correlation distance metric and repeat the above.

*Hint:*

---

```
1  distmethod <- function(x) as.dist(1-cor(x))  
2  distance <- distmethod(xTrain)
```

---

### 3 Species distribution modelling

<b>Data:</b>	<i>Bradypus variegatus</i> (brown-throated sloth) geographic distribution data ( <a href="#">Phillips et al. (2006)</a> )
<b>Task:</b>	produce a map showing the likely geographic distribution of <i>Bradypus variegatus</i> (unsupervised)
<b>Method:</b>	Gaussian mixture models

An ubiquitous challenge in conservation biology is to produce species distribution maps. The objective is to show how in different regions there is more or less chance of observing the species under consideration. The dataset consists of 116 latitude and longitude recordings of where *Bradypus variegatus* was observed. This is a classic density estimation problem which in this case is equivalent to a 2D histogram. Let us start by retrieving the dataset which is found in the `dismo` package and plot a point for every observation.

```
1 library(raster) # functions for gridded spatial data
2 library(dismo) # package containing the Bradypus variegatus dataset
3 library(rworldmap) # access to map of the world
4 dataPath <- file.path(system.file(package="dismo"), "ex") # path to data location
5 bradypusFilePath <- file.path(dataPath, "bradypus.csv") # path to Bradypus dataset
6 data <- read.table(file=bradypusFilePath, header=T, sep=",", skip=0) # read data
7 worldMap <- getMap(resolution = "low") # access world map, xlim = long, ylim = lat
8 plot(worldMap, xlim = c(-85, -40), ylim = c(-25, 20), asp = 1, axes=T)
9 points(data[, 2], data[, 3], pch=20, col="blue", cex=0.5) # add observed data points
```

The question now is which variables should we use to fit this histogram. A naïve approach is to use the observed latitude and longitude, after all we would expect to find the same species within short geographical distances from where we have already observed some. This approach however ignores the habitat properties inhabited by *Bradypus variegatus*. A better method would be to use environmental variables (e.g temperature and precipitation), which characterise the species' habitat, to fit this distribution, and then convert back to the latitude/longitude domain. The `dismo` package contains a number of bioclimatic variables stored as GRD files. This is a popular file format used in GIS (Geographic Information System) and contains grid cell values for that variable. The naming convention for these variables can be found [here](#)<sup>2</sup>. The data resolution is  $0.5^\circ \times 0.5^\circ$  which is sufficient for illustrative purposes. The steps taken in this task look something like this:

---

<sup>2</sup>Note that temperature data is stored as  $^\circ\text{C} \times 10$  as discussed [here](#)

1. Read in and plot environmental variables which cover the geographical area of interest.

```
1   grdFiles <- list.files(path=dataPath, pattern='grd', full.names=T) # read path of
   all .grd files
2   envData <- stack(grdFiles) # stack all environmental variables
3   plot(envData) # plot to confirm data is OK
```

2. For every observed data point, i.e 116 lat/long positions, extract a value for the environmental variable.

i.e (temperature, precipitation...)= $f$ (latitude, longitude)

```
1   xTrain <- extract(envData, data[, 2:3]) # columns 2 (longitude) and 3 (latitude)
```

3. Standardise the environmental data as ranges vary widely.

i.e  $x_{\text{standardise}} = \frac{x - \mu}{\sigma}$

This is a very important step and if omitted the algorithm might fail to converge

```
1   meanXTrain <- apply(xTrain, 2, mean) # compute mean of every column
2   sdXTrain <- apply(xTrain, 2, sd) # compute sd of every column
3   xTrain <- sweep(xTrain, 2, meanXTrain, FUN="-") # x - mean
4   xTrain <- sweep(xTrain, 2, sdXTrain, FUN="/") # (x - mean)/sigma
```

4. Fit a distribution to the standardised data using Gaussian mixture models. Instead of using all 9 variables however, only a subset of 3 are considered. The reasons are twofold: i) some variables are highly correlated with each other (e.g mean annual temperature and maximum temperature of warmest month), but most importantly because of ii) the *curse of dimensionality*. The more variables (dimensions) are considered the more the space will be sparsely populated, unless dataset is extremely large.

```
1   library(mclust)
2   envVariables <- c("bio1", "bio12", "biome") # chosen environmental variables
3   fit <- densityMclust(data=xTrain[, envVariables]) # fit distribution
```

5. Compute the value of the fitted distribution for every geographical grid of interest and plot the species geographical distribution map.

```
1   # Standardise test data (xTest is the whole grid not just the 116 obs.)
2   xTest <- as.data.frame(subset(envData, subset=envVariables))
3   xTest <- sweep(xTest, 2, meanXTrain[envVariables], FUN="-")
4   xTest <- sweep(xTest, 2, sdXTrain[envVariables], FUN="/")
5   # Read the distribution at the test points
6   notNA <- complete.cases(xTest) # ignore sea points (NA = sea)
7   pred <- rep(NA, dim(xTest)[1])
8   pred[notNA] <- predict(fit, newdata=xTest[notNA, ])
9   # Create a data frame with values for the distribution for every grid cell
10  df <- data.frame(pred=as.matrix(pred, nrow=nrow(envData), ncol=ncol(envData)))
11  coordinates(df) <- coordinates(envData) # set spatial coordinates
12  gridded(df) <- TRUE # coerce to SpatialPixelsDataFrame
13  rasterDF <- raster(df) # coerce to raster
14  library(RColorBrewer) # package for pretty colour palettes
15  plot(rasterDF, col=brewer.pal(n=9, name="Reds"))
16  points(data[, 2], data[, 3], pch=20, col="blue", cex=0.5) # add observations
```

**Question(s):**

1. Compare your geographical distribution map to [Phillips et al. \(2006\) pg. 252](#). Do you think that a Gaussian mixture model is appropriate for this problem?
2. Pick different combinations of **envVariables** and describe how the species distribution map changes. Try also single environmental variables.



## 4 Predicting forest cover type from cartographic attributes

<b>Data:</b>	forest cover type ( <a href="#">UC Irvine Machine Learning Repository</a> )
<b>Task:</b>	predict forest cover type for a set of cartographic measurements (supervised)
<b>Method:</b>	decision trees and random forests

This problem appeared as a competition on [Kaggle](#). It is a multi-class problem where several cartographic attributes are used to predict the forest cover type. Data comes from the Roosevelt National Forest of northern Colorado and each row represents an observation over a  $30\text{m} \times 30\text{m}$  patch. The seven distinct forest cover types (**Cover Type**) are:

1. Spruce/Fir
2. Lodgepole Pine
3. Ponderosa Pine
4. Cottonwood/Willow
5. Aspen
6. Douglas-fir
7. Krummholz

The cartographic attributes are:

**Elevation**

Elevation in meters

**Aspect**

Aspect in degrees azimuth

**Slope**

Slope in degrees

**Horizontal Distance To Hydrology**

Horizontal distance to nearest surface water features

**Vertical Distance To Hydrology**

Vertical distance to nearest surface water features

**Horizontal Distance To Roadways**

Horizontal distance to nearest roadway

**Hillshade 9am<sup>3</sup>**

Hillshade index at 9am, summer solstice

**Hillshade Noon**

Hillshade index at noon, summer solstice

---

<sup>3</sup>Hillshade index takes an integer value from 0 (complete shadow) to 255 (full sunlight)

Hillshade 3pm

Hillshade index at 3pm, summer solstice

Horizontal Distance To Fire Points

Horizontal distance to nearest wildfire ignition points

Wilderness Area

4 binary columns (0 = absence or 1 = presence) of wilderness area designation

Soil Type

40 binary columns (0 = absence or 1 = presence) of soil type designation

For information on wilderness area designation and soil type please check [additional information](#), however it is not essential for the purpose of this exercise. Let us start by reading the dataset<sup>4</sup> in and split it into a training and testing dataset.

```
1 library(RCurl) # To compose general HTTP requests
2 options(RCurlOptions = list(capath = system.file("CurlSSL", "cacert.pem", package = "
  RCurl"), ssl.verifypeer = FALSE)) # To avoid SSL certificate problem
3 dataPath <- getURL("https://dl.dropboxusercontent.com/u/57002389/ML_Life_Sciences/Data/
  ForestCoverData.csv") # Location of dataset
4 data <- read.table(text=dataPath, header=T, sep=",", skip=0) # Use a text connection
5 data <- data[complete.cases(data), ] # Only 2 cases of missing data so ignore
6 data$Cover_Type <- factor(data$Cover_Type)
7 NTOTAL <- dim(data)[1] # Total number of observations
8 NTRAIN <- 12000 # Let us use this many training datapoints
9 set.seed(101) # Just so we can reproduce results
10 trainIndices <- sample(x=NTOTAL, size=NTRAIN, replace=FALSE)
```

Fit a simple decision tree to the data.

```
1 library(rpart)
2 library(rpart.plot)
3 library(rattle) # For fancyRpartPlot to plot rpart tree nicely
4 fit <- rpart(Cover_Type ~ ., data=data, subset=trainIndices, method="class")
5 fancyRpartPlot(fit) # Plot decision tree
6 predClass <- predict(fit, type="class") # Pred class on training dataset
7 confusionMatrix <- table(data$Cover_Type[trainIndices], predClass)
8 classError <- 1 - diag(prop.table(confusionMatrix, 1)) # Compute misclassification rates
```

Results not promising, especially for Cover Type 4 to 7. Let us use a Random Forest and see how the classError changes.

```
1 library(randomForest)
2 fit <- randomForest(Cover_Type ~ ., data=data, subset=trainIndices, ntree=200,
  importance=TRUE) # Takes a few secs
3 fit$confusion # Confusion matrix
```

---

<sup>4</sup>For convenience I have uploaded it to Dropbox where it can be easily accessed using the RCurl package

The misclassification rate has improved, but for some **Cover Type** it is still unsatisfactory, why? A quick look at the confusion matrix reveals how our training dataset is *unbalanced*. **Cover Type** 1 and 2 comprise 85% of the training examples. Having one class more prevalent in a dataset is a common problem in classification, for example, you may have lots of data for healthy people but only a few for those suffering from a particular disease. There are two approaches to circumvent such problem:

1. **On a data level:** Data is balanced by under/over-sampling the original dataset. The aim is to get roughly equal number of classes in the training dataset
2. **On an algorithm level:** This depends on the algorithm used, but usually involves adjusting some objective function or specifying directly the class prior

**Question(s):**

1. Choose a training dataset that have roughly equal number of classes and repeat the fitting process (both a simple decision tree and Random Forest). What is the model's performance now? (compute misclassification error rates for both *training* and *testing* datasets)

*Hint:* `predClass <- predict(fit, newdata=data[-trainIndices, ])`

2. Repeat the above, but instead of balancing the dataset, specify the class prior by setting the `classwt` in the `randomForest` function. How does the model's performance compare with the above?

3. Using the model which gave you the best performance, can you elucidate on the most important variables?

*Hint:* Plot the variable importance plot `varImpPlot(fit)`

4. Instead of having 40 binary columns for **Soil Type** how else can you “code” this variable?

## 5 Classifying patients by their gene expression signature

**Data:** gene expression (data of T- and B-cell Acute Lymphocytic Leukemia)  
**Task:** predict a patient's phenotype using their gene expression signature (supervised)  
**Method:** decision trees

Here we will consider again the gene expression data from the previous clustering exercise. This time however we formulate it as a supervised learning problem. Given a gene expression signature can we predict the patient's phenotype? Let us start by considering only a binary classifier, that is, discriminating between B- and T-cell cancer types. As this is an illustrative example and the sample size is fairly small let us use all available data for training<sup>5</sup>.

```
1 library(ALL) # Load the ALL package
2 data(ALL) # Loads ALL dataset to workspace
3 xTrain <- t(exprs(ALL)) # Extract the 128 x 12625 dataset
4 yTrain <- factor(substr(ALL$BT,1,1)) # Class is either B or T
5 df <- data.frame(x=xTrain, y=yTrain) # Create a data frame, x - all inputs, y - class
```

### Question(s):

1. Try logistic regression, a very popular statistical binary classifier  
*Hint:* `glm(y ~ ., data=df, family=binomial(link="logit"))`  
What's the outcome and why?
2. Use decision trees now  
*Hint:* use the same code as listed in the previous Forest Cover Type exercise  
What's the classifier's performance on the training dataset?
3. Repeat the above but now we want to infer the stage of the disease as well  
*Hint:* Set `yTrain <-factor(ALL$BT)`

---

<sup>5</sup>Note that if we do this, we cannot infer the generalisation error. In this case using cross-validation would have been a more sensible option, but this was not discussed in detail in the workshop

## 6 Identifying patients with Parkinson's disease from speech recordings

<b>Data:</b>	audio recordings ( <a href="#">UC Irvine Machine Learning Repository</a> )
<b>Task:</b>	predict presence of Parkinson's disease (supervised)
<b>Method:</b>	support vector machine (SVM)

Parkinson's disease is a neurological condition that affects millions of people and their families. Around 90% of people with Parkinson's show some form of speech impairment. The premise of this study is to use audio recordings from a sustained phonations experiment<sup>6</sup> to classify patients as having Parkinson's or not; a binary classification problem. A total of 22 features (predictors), metrics related to the frequency content of the audio signal, were extracted from 195 recordings<sup>7</sup>. Let us load the dataset and scale all 22 features (feature scaling) to be in the -1 to +1 range ( $x' = \frac{2(x-x_{\min})}{(x_{\max}-x_{\min})} - 1$ ). The variable `status` is 0 or 1 indicating absence or presence of disease.

```
1 library(RCurl) # To compose general HTTP requests
2 options(RCurlOptions = list(capath = system.file("CurlSSL", "cacert.pem", package = "
  RCurl"), ssl.verifypeer = FALSE)) # To avoid SSL certificate problem
3 dataPath <- getURL("https://archive.ics.uci.edu/ml/machine-learning-databases/parkinsons
  /parkinsons.data") # Location of dataset
4 data <- read.table(text=dataPath, header=T, sep=",", skip=0, row.names=1) # Text read
5 data$status <- factor(data$status) # Convert to factor (0/1) (disease absence/presence)
6 iiPredictors <- which(names(data)!="status") # Column number for all (22) predictors
7 # Scale all features to -1 to +1 (all data columns except "status")
8 # x' = 2*(x - xmin)/(xmax - xmin) - 1
9 xMin <- apply(data[, iiPredictors], 2, min)
10 xMax <- apply(data[, iiPredictors], 2, max)
11 # Feature scaling
12 data[, iiPredictors] <- sweep(data[, iiPredictors], 2, xMin, FUN="-")
13 data[, iiPredictors] <- 2*data[, iiPredictors]
14 data[, iiPredictors] <- sweep(data[, iiPredictors], 2, (xMax-xMin), FUN="/")
15 data[, iiPredictors] <- data[, iiPredictors] - 1
```

Next we will split the data into a training (60%) and testing (40%) set. The testing dataset will be used to estimate the generalisation performance of the model. Apart from the misclassification rates and because we are dealing with a binary classification problem, we will also display the results as a receiver operating characteristic (ROC) curve. The various definitions (true positive, true negative etc...) are depicted in Figure 1.

---

<sup>6</sup>Patient is asked to produce and sustain a single vowel at constant pitch

<sup>7</sup>For further information on how these metrics were computed please refer to [Little et al. \(2009\)](#) but it is not essential for the purpose of this exercise

```

1 #-----#
2 # Split into training/testing 60% - 40%
3 set.seed(101) # Just so we can reproduce results
4 trainIndices <- sample(x=dim(data)[1], size=round(dim(data)[1]*0.6), replace=FALSE)
5 #-----#
6 # Train SVM
7 library(e1071) # A support vector machine (SVM) library (there is also kernlab)
8 fit <- svm(status ~., data=data, subset=trainIndices, type="C-classification", kernel="
   linear", probability=TRUE)
9 #-----#
10 # Compute misclassification rates on testing dataset
11 predClass <- predict(fit, newdata=data[-trainIndices, ], type="class") # 0/1
12 confusionMatrix <- table(data$status[-trainIndices], predClass)
13 classError <- 1 - diag(prop.table(confusionMatrix, 1))
14 #-----#
15 # Plot Receiver operating characteristic (ROC)
16 library(ROCR) # ROC curves library
17 # Compute prediction probabilities rather than just class
18 predProb <- predict(fit, newdata=data[-trainIndices, ], probability=TRUE)
19 # Prediction object needed by ROCR
20 # prediction(probability of class=positive class, actual class label)
21 predObj <- prediction(attr(predProb, "probabilities")[, 1], data$status[-trainIndices])
22 ROCObj <- performance(predObj, "tpr", "fpr") # ROC object
23 plot(ROCObj) # Plot false positive rate vs true positive rate
24 abline(a=0, b=1, col="red", lty=2) # Randomly guessing line
25 #-----#
26 # Compute area under ROC curve (AUC)
27 # AUC = 0.5 (randomly guessing class), AUC = 1 (perfect predictor)
28 areaUnderCurve <- performance(predObj, "auc") # Compute AUC
29 areaUnderCurve <- unlist(slot(areaUnderCurve, "y.values")) # Extract AUC (converting S4
   class to vector)

```




Figure 1: Binary classification definitions for a hypothetical predictor with a class cut-off depicted by the dotted black line

**Question(s):**

1. Why is the misclassification rate of the “no disease” class much worse than the “disease” class? What can we do?
2. Can we reduce the number of predictors *without* taking into consideration the output class (unsupervised feature selection)?

*Hint:* In this case, because all the predictors for an observation were derived from the *same* audio recording, we can deduce that some metrics are highly correlated and thus introducing no further knowledge to our dataset

---

```
1 library(corrgram)
2 corrgram(data[, iiPredictors], lower.panel=panel.shade, upper.panel=panel.conf)
```

---

3. Repeat the analysis and compare results when using the reduced feature set, and also for the case of a radial basis function kernel

*Hint:* Use `kernel="radial"`